

FMAP: the flow sequence mapping program ¹

Nils Homer

¹Copyright ©2010 Ion Torrent Systems, Inc. All Rights, made for FMAP version 0.0.1

Contents

Table of Contents	iii
Preface	v
1 Basic Usage	1
1.1 Program Organization	1
1.2 Common Mapping Options	1
1.2.1 Global Options	2
1.3 fmap index	4
1.3.1 Usage	4
1.4 fmap server	5
1.4.1 Usage	5
1.5 fmap map1	6
1.5.1 Usage	6
1.6 fmap map2	7
1.6.1 Usage	8
1.7 fmap map3	8
1.7.1 Usage	9
1.8 fmap mapall	9
2 File Formats	11
2.1 SAM Alignment Format	11
2.1.1 Optional Tags	11
3 Appendix	13
References	13

Preface

This document is meant to serve as a guide for the practical use of FMAP. It includes explanations of all command-line options for each command and binary in FMAP to give an idea of basic usage. Input and output file formats are also detailed. The default options of the various programs are designed intelligently to adapt to the various program uses, but in some cases the options may need to be customized.

This document does not try to explain the underlying algorithm or data-structures used in FMAP. Without proper understanding of the underlying algorithms, it is difficult to use this very flexible program knowledgeably to obtain your desired results.

If you have anything that you would be useful to add to this guide, feel free to relay the addition to the FMAP developers. This includes but is not limited to bugs, typos, and explanations. Please see <http://www.iontorrent.com> for more details.

Enjoy!

Chapter 1

Basic Usage

1.1 Program Organization

FMAP consists of a set of utilities, combined into one command line program called `fmap`. Each utility is specified by a unique name or command. Specifying FMAP without a command will give a list of commands, while specifying FMAP with a command but no options will give a list of options for that command.

There are two steps in mapping with FMAP. The first step is to build an index of the reference genome onto which we map. This index needs only to be built once for each genome, performing much of the mapping work upfront. The second step is to map the reads to the reference genome using this index.

FMAP is implemented as a command-line program. It accepts many command-line options to customize and tune the mapping algorithm. The key commands are organized into one binary program called `fmap`. To access each command, we use `fmap <name>`, where `<name>` is the name of the command we wish to execute.

1.2 Common Mapping Options

Some common options exist across some or all of the mapping commands (`map1`, `map2`, `map3` and `mapall`). These options will be discussed here to avoid duplication.

The FMAP mapping commands can accept their input file from the standard input stream. The output of these commands also write to the standard output stream. This facilitates the use of these FMAP commands in a pipe-and-filter model.

1.2.1 Global Options

`-f FILE`

Specifies the file name of the reference genome in FASTA format. The maximum number of bases in the reference genome can be 4294967295, due to the use of 32-bit values to represent each position in the reference.

`-r FILE`

Specifies the file name of the reads to mapped. The input reads can be in FASTA, FASTQ, or SFF format.

`-F STRING`

Specifies the file format of the reads file. Without this option, the format is auto-recognized by the file extension. The valid inputs are `fa` or `fasta` for FASTA, `fq` or `fastq` for FASTQ, or `sff` for SFF.

`-A INT`

Specifies the match score. This number must always be positive.

`-M INT`

Specifies the mismatch penalty. This number must always be positive.

`-O INT`

Specifies the gap open penalty. This number must always be positive.

`-E INT`

Specifies the gap extension penalty. This number must always be positive.

`-w INT`

Specifies the band width for local alignment. This number must always be positive.

`-g`

Specifies that the entire read should be mapped to the reference. Otherwise a sub-sequence of the read may be mapped.

`-q INT`

Specifies the number of reads to cache or load into memory at one time.

`-n INT`

Specifies the number of threads to run.

`-a INT`

Specifies the output filter for the mappings.

0. returns the mapping with the best score only if all other mappings had worse score, otherwise the read is unmapped.
1. returns the mapping with the best score. If more than one mapping has this score, a random mapping with this score is returned.
2. returns all the mappings with the best score.
3. returns all the mappings, regardless of score.

Reads that have no mapping are returned as unmapped reads.

`-R STRING`

Specifies the RG (read group) line to use in the SAM file. The “PG”, “FO”, and “KS” tags should not be specified in this string; they will be added by fmap.

`-j, -z`

Specifies that the input is bzip2 (`-j`) or gzip (`-z`) compressed. This is auto-recognized if the input file name has the extension `.bz2` for bzip2 and `.gz` for gzip.

`-J, -Z`

Specifies that the output should be bzip2 (`-J`) or gzip (`-Z`) compressed.

`-s INT`

Specifies the shared memory key if the reference index has been loaded into shared memory.

`-v`

Specifies to print verbose progress messages, otherwise progress messages will be suppressed.

`-h`

Specifies to print a help message, listing all available options.

1.3 fmap index

The **index** command creates a compact version of the reference genome and associated index. The index is stored as a compressed suffix array using the FM-index and BWT transform (Ferragina and Manzini (2000); Burrows and Wheeler (1994)). A hash into this index accelerates this lookups of DNA sequences in this index. In fact, a second additional index of the reference genome is created that indexes the reverse (but not complimented) reference genome. This second index further speeds up the search time.

See [section 1.2](#) for common options that are in use in this command.

1.3.1 Usage

`-o INT`

Specifies the occurrence interval size o , storing only every o th occurrence interval. This must be one, or a multiple of 16.

`-w INT`

Specifies the k-mer size (the number of bases) to hash. The size of the hash give the k-mer size k in bytes is:

$$= 2 \sum_{n=1}^k 4^n = 2 \left(\frac{1 - 4^{k+1}}{1 - 4} - 1 \right)$$

using the Taylor series.

`-i INT`

Specifies the suffix array interval size *i*, storing only every *i*th suffix array interval. This must be one, or a multiple of two.

`-a STRING`

Specifies the BWT construction algorithm (`bwtsw` or `is`). The `bwtsw` algorithm is for genomes larger than or equal to 10Mb, and the `is` algorithm is for genomes smaller than 10Mb. This will be auto-recognized during index creation if this option is omitted.

1.4 fmap server

The **server** command loads the reference genome index and data into shared memory. This lets other mapping instances avoid having to load the index upon each execution. This program will try to fail gracefully, detaching shared memory upon exiting.

See [section 1.2](#) for common options that are in use in this command.

1.4.1 Usage

`-c STRING`

Specifies the command for the server (`start`, `stop`, `kill`). The `start` command loads data into shared memory, and waits for a `ctrl-c` or `SIGINT` signal. The `stop` command will signal a currently running server to stop and detach shared memory. The `kill` command will forcibly detach the shared memory segment and signal a currently running server to exit. The latter command is especially useful for killing zombied processes and detaching lost shared memory segments.

`-k INT`

Specifies the shared memory key for this server. The shared memory segment will be identified by this key.

`-a`

Specifies to load all reference genome data structures into memory.

-r

Specifies to load the forward packed reference sequence.

-R

Specifies to load the reverse packed reference sequence.

-b

Specifies to load the forward BWT sequence.

-B

Specifies to load the reverse BWT sequence.

-s

Specifies to load the forward suffix array.

-S

Specifies to load the reverse suffix array.

1.5 fmap map1

The **map1** is a command to quickly map short sequences to a reference genome by intelligently enumerating errors. This algorithm is not well suited for longer reads ($< 150\text{bp}$) and based off of the BWA short-read algorithm ([Li and Durbin \(2010a\)](#)).

See [section 1.2](#) for common options that are in use in this command.

1.5.1 Usage

-l INT

Specifies the k-mer length for seeding candidate alignment locations (CALs) using the first k bases of the read.

-k INT

Specifies the number of mismatches allowed in the seed.

-m NUM

Specifies the maximum number of mismatches allowed in the mapping, or the fraction of mismatches with respect to the read length.

-o NUM

Specifies the maximum number of indels allowed in the mapping, or the fraction of indels with respect to the read length. The number of indels is equal to the number of gap opens.

-e NUM

Specifies the maximum number of indel extensions allowed in the mapping, or the fraction of indel extensions with respect to the read length. The number of indel extensions is equal to the number of gap extensions.

-d INT

Specifies the maximum number of candidate alignment locations (CALs) to allow a deletion to be extended.

-i INT

Specifies to disallow indels within this number of bases from the ends of the read.

-b INT

Specifies to stop searching for mappings when this many best scoring mappings have been found.

-Q INT

Specifies the maximum number of alignment nodes in memory in the implicate prefix trie traversal before the search is stopped.

1.6 fmap map2

The **map2** is a command to quickly map long sequences to a reference genome. This algorithm is well suited for longer reads (≥ 150 bp) and based off of the BWA long-read algorithm (Li and Durbin (2010b)).

See [section 1.2](#) for common options that are in use in this command.

1.6.1 Usage

-c FLOAT

Specifies the coefficient to adjust the mapping score threshold based on the read length. Given a l -long read, the threshold for a mapping to be retained is $a * \max\{T, c * \log(l)\}$, where a is the match score ($-A$) and T is the minimum score threshold ($-T$).

-w INT

Specifies the band width in the banded local alignment.

-T INT

Specifies the number of multiples of the match score ($-A$) for the minimum scoring threshold.

-S INT

Specifies the maximum seeding interval size for extending a mapping.

-b INT

Specifies the number of mappings to retain at during extension of the seed.

-N INT

Specifies the number seeds supporting the given mapping needed to skip the reverse alignment.

1.7 fmap map3

The **map3** is a command to map sequences to a reference genome. This algorithm is well suited for longer reads (≥ 150 bp) and a simplification of the SSAHA long-read algorithm (Ning *et al.* (2001)).

See [section 1.2](#) for common options that are in use in this command.

1.7.1 Usage

-l INT

Specifies the k -mer length to seed candidate alignment locations (CALs). With a value of -1 will set the seed length to $(\text{ceiling}(\log_4(R)) + 2)$ where R is the reference genome length.

-S INT

Specifies the maximum number of CALs allowed to be returned by a seed before it is ignored.

-b INT

Specifies the window in bases in which to group seeds.

-w INT

Specifies the band width in the banded local alignment.

-T INT

Specifies the number of multiples of the match score ($-A$) for the minimum scoring threshold.

1.8 fmap mapall

The **mapall** is a command to quickly map short sequences to a reference genome. This command combines available mapping algorithms for fast and sensitive alignment. The algorithm follows a two-stage approach, with a set of algorithms and associated settings for each stage. If there are no mappings for a read by applying the algorithms in the first stage, then the algorithms in the second stage are applied. For example, a set of algorithms to quickly align near-perfect reads may be used in the first stage, while a set of sensitive algorithms may be used to map difficult reads in the second stage.

Note: for the `map1` command used within `mapall`, the alignment score is adjusted to agree with alignment scores from other algorithms. This may cause the alignment score reported by using `map1` within the `mapall` command to be different than using `map1` by itself.

See [section 1.2](#) for common options that are in use in this command.

Usage

First, a set of global options is specified that will be used for the algorithms to be applied at both stages. Global options should not be given in the options for a specific algorithm. Next, the algorithms and their associated options will be specified for each stage. An algorithm to be applied in the first stage should be specified using its name in lowercase. The options specific to this algorithm should be specified directly thereafter. An algorithm to be applied in the second stage should be specified using its name in uppercase. The options specific to this algorithm should be specified directly thereafter.

An example would be:

```
fmap mapall -f ref.fasta -r reads.fastq -g -M 3 map1 -l 12 -k 24 MAP2  
-b 5 MAP3 -S 10.
```

In this case, the map1 algorithm will be applied with the options `-l 12 -k 24 -M 3` in the first stage. If no mapping is found for a read, the map2 algorithm with the options `-b 5 -g -M 3` and the map3 `-g -M 3` algorithm with the options `-S 10` will be applied in the second stage. Notice how the global options `-g -M 3` are applied to all the algorithms were applicable. It is possible to have a given algorithm run in both stages with different options.

Chapter 2

File Formats

2.1 SAM Alignment Format

FMAP produces mappings in the SAM format ([Li *et al.* \(2009\)](#)). Optional tags are used to store information about mappings useful for downstream processing.

2.1.1 Optional Tags

AS

This stores the alignment score.

NM

This stores the edit distance from the reference sequence.

XM

This stores the number of mismatches in the mapping.

XO

This stores the number of indels (gap opens) in the mapping.

XG

This stores the number of indel extensions (gap extensions) in the mapping.

XS

This stores the alignment score of next-best sub-optimal mapping.

XF

This stores from where the mappings originated: one indicates from the forward search, two indicates from the reverse search, and three indicates from both the forward and reverse search.

XE

This stores the number of seeds supporting this mapping.

XI

This stores the size of the suffix interval for this mapping.

XA

This stores the algorithm that produced this mapping and from what stage. The format is the algorithm name, and then the zero-based stage, separated by a dash.

Chapter 3

Appendix

Bibliography

- Burrows, M. and Wheeler, D. (1994). A block sorting lossless data compression algorithm. Technical Report 124, Digital Equipment Corporation.
- Ferragina, P. and Manzini, G. (2000). Opportunistic data structures with applications. In *Proceedings of the 41st Annual Symposium on Foundations of Computer Science*, pages 390–, Washington, DC, USA. IEEE Computer Society.
- Li, H. and Durbin, R. (2010a). Fast and accurate long-read alignment with Burrows-Wheeler transform. *Bioinformatics*, **26**, 589–595.
- Li, H. and Durbin, R. (2010b). Fast and accurate long-read alignment with Burrows-Wheeler transform. *Bioinformatics*, **26**, 589–595.
- Li, H., Handsaker, B., Wysoker, A., Fennell, T., Ruan, J., Homer, N., Marth, G., Abecasis, G., and Durbin, R. (2009). The Sequence Alignment/Map format and SAMtools. *Bioinformatics*, **25**, 2078–2079.
- Ning, Z., Cox, A., and Mullikin, J. (2001). SSAHA: a fast search method for large DNA databases. *Genome Res.*, **11**, 1725–1729.